

L10: Machine learning use cases in algorithmic trading

Course Code: COMP7415

Agenda

- Major categories of machine learning algorithms
- Multiple use cases of machine learning in algo-trading
- The future of smart investment systems

Introduction to Machine Learning

What is Machine Learning?


- A subset of artificial intelligence (AI) that enables systems to learn from data, identify patterns, and make decisions with minimal human intervention.
- Key Characteristics:
 - Data-driven
 - Adaptive learning

Types of Machine Learning

1. Supervised Learning
2. Unsupervised Learning
3. Reinforcement Learning

1. Supervised Learning

- Involves training a model on a labelled dataset where the outcome is known.
- Two main category of algorithms
 1. Classification
 - predict categorical target variables
 - Eg. SVM, Logistic Regression, Decision Tree, CNN, etc
 2. Regression
 - predict continuous target
 - Eg. Linear regression, decision tree, random forest, etc



1. Supervised Learning

Pros:



- Typically delivers high accuracy if the training data is representative and of good quality.
- The presence of labelled data provides a clear objective for the model, making it easier to evaluate performance.
- Model Interpretability: Many supervised learning algorithms offer insights into feature importance and decision-making processes.

Cons:

- Requires a large amount of labelled data, which can be expensive and time-consuming to obtain.
- The model can only predict outcomes it has been trained on, making it ineffective for novel or unrepresented scenarios.

2. Unsupervised Learning

- Discover patterns and relationship in data without labelling
- Two main category of algorithms
 1. Clustering
 - group data points into clusters based on their similarity
 - Eg. K-mean clustering, PCA, etc
 2. Association
 - identifies rules that indicate the presence of one item implies the presence of another item
 - Eg. Apriori algorithm, Eclat algorithm, etc




2. Unsupervised Learning

- Pros:
 - Can work with unlabelled data, making it easier to gather and analyze large datasets without the need for extensive labelling.
 - Capable of uncovering complex structures and relationships in data that may not be apparent, leading to new insights.
- Cons:
 - Results can be difficult to interpret, as there are no clear targets or labels to guide understanding of the outcomes.
 - May not yield as high predictive accuracy as supervised learning methods, especially for specific tasks like classification.
 - Without labelled data, there's a risk of identifying patterns that are not meaningful or relevant, leading to incorrect conclusions.

3. Reinforcement Learning

- An agent interacts with the environment by producing actions and discovering errors, and learns to make decisions by maximizing cumulative reward.
- Common algorithms:
 - Q-learning
 - SARSA (State-Action-Reward-State-Action)



3. Reinforcement Learning

- Pros:
 - Learns optimal strategies through trial and error, minimizing the need for human intervention and supervision.
 - Can adapt to dynamic environments, continuously improving performance based on feedback and changing conditions.
 - Well-suited for complex problems with long-term decision-making, such as game playing and robotic control, where actions have delayed rewards.
- Cons:
 - Requires a large number of interactions with the environment to learn effectively, which can be time-consuming and resource-intensive.
 - Define appropriate reward functions can be challenging and may lead to unintended behaviors if not designed carefully.

Summary

	Description	Data Requirement	Common Use Cases
Supervised Learning	Learns from labelled data to make predictions.	Labeled data is required.	Classification, regression tasks.
Unsupervised Learning	Finds patterns or structures in unlabelled data.	No labelled data needed.	Clustering, dimensionality reduction.
Reinforcement Learning	Learns through trial and error by receiving feedback.	Interaction with environment.	Game playing, robotics, autonomous systems.

Supervised Learning in Algo-Trading

What is a Trading Strategy?

$$y = f(x)$$

Outputs:

- action (i.e. buy/ sell/ do nothing)
- quantity
- price
- stop loss
- take profit
- ...



Model / Logic:

- technical analysis
- fundamental analysis
- pricing model
- statistical arbitrage
- regression
- ...



Inputs:

- timeframe
- price
- volume
- order book
- technical indicators
- fundamental data
- news
- ...



RSI Example



Make sense?

Inputs:

- market price
- 14-day RSI value

Logic:

- buy if RSI < 30
- sell if RSI > 70
- do nothing otherwise


Output:

- buy/sell action




Traditional vs AI Approach

Rule driven: define a rule, then see what happens later.



Result driven: identify a desirable trade, then see what happened earlier.



RSI Example



100% accuracy?



Overfitting?

Output:

- desire to buy at red points
- desire to sell at yellow points

Inputs:

- market price
- any technical indicators

Logic:

- What should be the hidden rules?



Unsupervised Learning in Algo-Trading

Overview

- Unsupervised learning analyzes financial data without labelled outcomes
- Applications in Trading:
 - Clustering for Market Segmentation
 - Grouping similar stocks or assets based on historical price movements or features to identify trends and potential investment opportunities
 - Anomaly Detection
 - Identifying unusual market behaviors or price movements that may signal trading opportunities or risks, helping traders react to potential market shifts
 - Feature Extraction
 - Reducing dimensionality of financial datasets (eg. PCA) to highlight key factors influencing asset prices, aiding in more effective modeling

Stock Clustering

- Involves grouping stocks based on their historical price movements or other financial metrics, revealing similarities that can inform trading strategies.
- Use cases:
 - Identifying Investment Opportunities
 - Helps in discovering stocks that are moving together, indicating potential correlations
 - Further develop pair-trading/ co-integration strategies
 - Portfolio Diversification
 - Build diversified portfolios by selecting stocks from different clusters

K-mean Clustering


- **Objective:** Minimize the variance within each cluster while maximizing the variance between clusters.
- How it works
 1. Initialization: Randomly select K initial centroids from the dataset
 2. Assignment: Assign each data point to the nearest centroid based on Euclidean distance.
 3. Update: Recalculate the centroids as the mean of the assigned data points.
 4. Repeat: Iterate the assignment and update steps until convergence (i.e., centroids no longer change).

K-mean Clustering


1. Define the number of cluster K
2. Initialize K centroids randomly
3. Repeat until convergence:
 - a. For each data point, assign it to the nearest centroid
 - b. For each centroid, update the centroid to the mean of assigned points




(a)




(b)




(c)



(d)




(e)



(f)

K-means in Python



```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

# Generating synthetic data
np.random.seed(42)
X1 = np.random.randn(50, 2) + np.array([0, 0]) # Cluster 1
X2 = np.random.randn(50, 2) + np.array([5, 5]) # Cluster 2
X3 = np.random.randn(50, 2) + np.array([10, 0]) # Cluster 3

# Combine clusters into one dataset
X = np.vstack([X1, X2, X3])

# Applying K-Means
kmeans = KMeans(n_clusters=3, random_state=0).fit(X)

# Getting the cluster centers and labels
centroids = kmeans.cluster_centers_
labels = kmeans.labels_

# Plotting the results
plt.scatter(X[:, 0], X[:, 1], c=labels, s=100, cmap='viridis')
plt.scatter(centroids[:, 0], centroids[:, 1], c='red', s=200, alpha=0.75)
plt.title('K-Means Clustering')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```

K-means for stock cluster

- Group the following 7 stocks into 2 clusters, based on their historical returns in the year of 2024
 1. MSFT: Microsoft
 2. GOOG: Google
 3. AAPL: Apple
 4. TSLA: Tesla
 5. NVDA: Nvidia
 6. JPM: JP Morgan
 7. BAC: Bank of America

K-means for stock cluster

```
import pandas as pd
from sklearn.cluster import KMeans
import yfinance as yf

# Define the stock symbols
stocks = ['MSFT', 'GOOG', 'AAPL', 'TSLA', 'NVDA', 'JPM', 'BAC']


# Collect historical price data for the last year
data = yf.download(stocks, start='2024-01-01', end='2024-12-31')['Close']

# Calculate daily returns
returns = data.pct_change().dropna()

# Apply K-Means clustering
kmeans = KMeans(n_clusters=2, random_state=0)
kmeans.fit(returns.T) # Transpose to cluster stocks instead of time periods

# Get the cluster labels
labels = kmeans.labels_

# Create a DataFrame to hold the results
results = pd.DataFrame({'Stock': stocks, 'Cluster': labels})
print(results)
```




	Stock	Cluster
0	MSFT	0
1	GOOG	0
2	AAPL	0
3	TSLA	0
4	NVDA	0
5	JPM	0
6	BAC	1

K-means for stock cluster

- By default, scikit-learn uses Euclidean distance for calculating the distance between data points (in above case, the feature vectors representing the time series of return).
- For 2 time series,
 - stock A = $\{r_{A,1}, r_{A,2}, \dots, r_{A,T}\}$
 - stock B = $\{r_{B,1}, r_{B,2}, \dots, r_{B,T}\}$
- The Euclidean distance will be

$$d(A, B) = \sqrt{\sum_{i=1}^T (r_{A,i} - r_{B,i})^2}$$




Cosine Distance

- $\cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$
- $d(A,B) = 1 - \cos(\theta) = 1 - \frac{A \cdot B}{\|A\| \|B\|}$

where

- $\|A\| = \sqrt{\sum_{i=1}^T r_{A,i}^2}$
- $\|B\| = \sqrt{\sum_{i=1}^T r_{B,i}^2}$
- $A \cdot B = \sum_{i=1}^T r_{A,i} \times r_{B,i}$

- $\cos(\theta)$ is ranging from -1 to 1
- Thus $d(A,B)$ is ranging from 0 to 2
- If $d(A,B)$ is close to 0, the 2 datasets are very similar or close to each other.



Correlation-Based Distance

- Value is ranging from 0 to 2
- When the distance d is close to 0, the 2 datasets are very similar

$$d(A, B) = 1 - \text{corr}(A, B) = 1 - \frac{\text{Cov}(A, B)}{\sqrt{\text{Var}(A)\text{Var}(B)}}$$

Limitation of above distance measures

- Assumes that the time series are aligned and of equal length.
- If the time series have different lengths or require alignment (eg. due to different time frames), we need to preprocess the data accordingly.


Limitation of K-means

- The results is sensitive to the initial placement of centroids, potentially leading to different cluster assignments in different runs.
- The algorithm requires to specify the number of clusters (k) in advance, which can be challenging if the optimal number is unknown.
- K-Means is sensitive to outliers, as they can disproportionately influence the position of the centroids, skewing the clustering results.

Reinforcement Learning in Algo-Trading

Overview

- A type of machine learning where an agent learns to make decisions by taking actions in an environment to maximize cumulative rewards.
- Key components:
 - **Agent:** The trading algorithm.
 - **Environment:** The stock market.
 - **Actions:** (eg. buy, sell, hold)
 - **Rewards:** Profit or loss from trades



SCORE: 0

Q-Learning Algorithm

- A model-free reinforcement learning algorithm used to find the optimal action-selection policy for an agent interacting with an environment.
- Enable the agent to learn how to maximize cumulative rewards over time by estimating the value of taking certain actions in given states.
- Denote:
 - State (s): A representation of the current situation in the environment
 - Action (a): The choices available to the agent that can change the state.
 - Reward (r): A feedback signal received after taking an action in a particular state, indicating the immediate benefit of that action.
 - Q-Value ($Q(s, a)$): Represents the expected future rewards for taking action a in state s .
 - Policy (π): The agent's strategy in deciding actions based on the current state

Q-Learning Steps

1. **Initialization:** Start by initializing the Q-table with zeros for all state-action pairs.
2. **Exploration:** At each step, choose an action based on the current state, either by exploring (choosing a random action) or exploiting (choosing the action with the highest Q-value).
3. **Taking Action:** Execute the chosen action, observe the reward and the next state.
4. **Update Q-Value:** Use the Q-learning formula to update the Q-value for the state-action pair.
5. **Iterate:** Repeat the process for many episodes, allowing the agent to learn the optimal policy over time.

Q-Learning Formula


- Q-value update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$

where

- $Q(s, a)$: Current estimated value of taking action a in state s .
- α : Learning rate ($0 < \alpha \leq 1$), determines how much new information overrides the old information.
- r : Immediate reward received after taking action a from state s .
- γ : Discount factor ($0 \leq \gamma < 1$), which balances the importance of immediate and future rewards.
- s' : The new state after taking action a .
- $\pi(a) = \max_{a'} Q(s', a')$: The maximum predicted Q-value for the next state, indicating the best possible action from state s' .

Example




Heuristic Function

- A heuristic function is a technique used to improve the efficiency of a search algorithm.
- It provides an estimate of the cost to reach a goal state from a given state.
- Role in Q-Learning
 - In reinforcement learning, especially Q-learning, heuristic functions can guide the agent's exploration strategy.
 - Utilizing prior knowledge about the environment to prioritize actions that are more likely to lead to better rewards.

Heuristic Function Example

- In grid-based environments, estimating the distance to the goal can guide the agent's movements.
- In a pac-man game, we want the agent to
 - Move towards food
 - Avoid moving towards ghosts



Benefits of Using Heuristic Functions

- **Faster Convergence:** Heuristics can speed up the learning process by focusing on promising actions.
- **Improved Performance:** They can lead to higher cumulative rewards by avoiding suboptimal paths.
- **Efficient Exploration:** Heuristics reduce the amount of exploration needed by directing the agent towards beneficial states.

Apply Q-Learning to Trading

Define the basic structure

- **State Space:**
 - Current market conditions (eg. price, volume)
- **Action Space:**
 - Only 3 possible actions
 - 1: buy
 - -1: sell
 - 0: hold or do nothing
- **Reward Function:**
 - Profit or loss from an action

Python Example

- Define trading environment

```
import numpy as np
import pandas as pd
import random

class TradingEnvironment:
    def __init__(self, data):
        self.data = data
        self.current_step = 0
        self.balance = 1000 # Starting balance
        self.shares = 0

    def reset(self):
        self.current_step = 0
        self.balance = 1000
        self.shares = 0
        return self.current_step # Return the current step as the state index

    def step(self, action):
        daily_return = self.data.iloc[self.current_step]
        reward = 0
        if action == 1: # Buy
            self.shares += 1
            self.balance *= (1+daily_return)
        elif action == 2: # Sell
            # long only
            if self.shares > 0:
                self.shares -= 1
                self.balance *= (1- daily_return)

        reward = self.balance - 1000 # Total profit
        self.current_step += 1
        done = self.current_step >= len(self.data) # Check if done
        # terminal states
        if done:
            next_state = None # No next state if done # non-terminal states
        else:
            next_state = self.current_step # Current step as next state

        return next_state, reward, done # Return the next state
```

- Define q-learning agent

```
class QLearningAgent:
    def __init__(self, state_size, action_size):
        self.q_table = np.zeros((state_size, action_size))
        self.learning_rate = 0.1
        self.discount_factor = 0.95
        self.exploration_prob = 0.9

    def choose_action(self, state):
        if random.uniform(0, 1) < self.exploration_prob:
            return random.choice(range(len(self.q_table[0]))) # Explore
        else:
            return np.argmax(self.q_table[state]) # Exploit

    def learn(self, state, action, reward, next_state):
        if next_state is not None: # Only learn if next_state is not terminal
            best_next_action = np.argmax(self.q_table[next_state])
            td_target = reward + self.discount_factor * self.q_table[next_state][best_next_action]
            td_delta = td_target - self.q_table[state][action]
            self.q_table[state][action] += self.learning_rate * td_delta
```

- Iterate the game for 1000 times

```
# Sample price data
data = pd.Series([100, 101, 102, 99, 98, 100, 105])
daily_returns = (data - data.shift(1)) / data.shift(1)
daily_returns = daily_returns.dropna()


env = TradingEnvironment(daily_returns) # Define the number of states and actions

state_size = len(data) # Number of steps as states
action_size = 3 # Actions: Hold (0), Buy (1), Sell (2)
agent = QLearningAgent(state_size, action_size) # Training the agent

for episode in range(1000):
    state = env.reset()
    done = False
    while not done:
        action = agent.choose_action(state)
        next_state, reward, done = env.step(action)
        agent.learn(state, action, reward, next_state)
        state = next_state if next_state is not None else state # Stay in the same state if done

# Print the final policy
final_policy = np.argmax(agent.q_table, axis=1)
action_labels = {0: 'Hold', 1: 'Buy', 2: 'Sell'}

print("Final Policy:", final_policy)
for state, action in enumerate(final_policy):
    print(f"State {state}: {action_labels[action]}")
```

A terminal window with a white title bar and three colored window control buttons (red, yellow, green) on the left. The terminal displays the output of the Python script in green text on a black background. The output shows the final policy as a list of actions for each state, followed by a mapping of state numbers to action names.

```
Final Policy: [1 1 2 2 1 0 0]
State 0: Buy
State 1: Buy
State 2: Sell
State 3: Sell
State 4: Buy
State 5: Hold
State 6: Hold
```

Considerations

- How to define a market state?
 - Daily returns, volume, market trend, volatility, etc?
 - Unlike Go, financial market is a non-finite and continuous environment
- How to define the action space?
 - In previous Q-learning example, we set {buy, sell, do nothing}
 - It doesn't reflect all possible trading actions in reality (eg. place a limit order at different price levels, dynamic trading size, etc)
- Given the current market condition, there are many possibilities for the future. Which one is the most likely to happen?
 - Is it possible to apply Monte Carlo tree search to simulate future market states?
 - Can the trading agent self-learn from the simulated environment, like the self-playing AlphaGo?




Limitation of Q-Learning in Trading

- High Dimensionality & Computational Complexity
 - Trading environments often have large state and action spaces (eg. multiple assets, different strategies) which may require significant computational power and time to train Q-learning models.
- Non-Stationary Environment
 - Financial markets are influenced by numerous factors that change over time.
 - Q-learning assumes a stationary environment, which may lead to outdated policies as market conditions evolve.
- Overfitting Risks
 - Q-learning models may overfit to historical price data, performing well in backtesting but poorly in live trading.
 - Strategies that work in one market condition may not generalize to others.

The Future of Investment

BTC \$105,888.00	ETH \$3,546.05	SOL \$166.72	BNB \$984.00	DOGE \$0.1794	XRP \$2.54
---------------------	-------------------	-----------------	-----------------	------------------	---------------



Alpha Arena Season 1 is now over, as of Nov 3rd, 2025 5 p.m. EST

Season 1.5 coming soon

Emerging Technologies in Algo-Trading

- Sentiment Analysis & Natural Language Processing (NLP)
- Big Data Analytics
- AI and Machine Learning
- Explainable AI (XAI)
- Blockchain
- Cloud & Edge Computing
- Quantum computing

A smart investment system should be able to

- Predict which direction the markets and assets will move
- Detect the potential risks and their impact
- Choose the best timing to trade
- Optimize for transaction costs
- Manage the position, risk and capital
- Explain the decisions it made
- Generate new investment strategies
- Know what other investors think and predict the actions they will take
- ...

Questions

1. Will AI or investment bot beat the best fund manager?
2. Given that you have access to all data in the world (both real-time and historical, public and private), are you able to accurately predict a stock price 10 minutes later?
3. In a world dominated by trading algorithms,
 - What will be the role of human traders?
 - Will the market become more or less volatile?

Key Takeaways

- Understand different types of machine learning algorithms
- Learn multiple use cases of machine learning in algo-trading
 - Data labelling on desired trades
 - Stock segmentation
 - Q-learning
- Understand the limitation of AI, so synthesis of human and AI is important